

AI Documentation

Module-wise AI Features, Capabilities & Configuration Guide

Table of Contents

1. AI Platform Overview	3
1.1 z-ai-web-dev-sdk Integration	3
1.2 AI Capabilities Across the Platform	3
2. AI Admin Console	4
2.1 Model Configuration	4
2.2 Prompt Templates	4
2.3 Monitoring Dashboard	5
2.4 Cost Tracking & Budget Management	6
3. AI Interview Module	6
3.1 Question Generation	6
3.2 Automated Assessment	7
3.3 Scoring Algorithms	8
3.4 Feedback Generation	8
3.5 Bias Detection	9
3.6 Interview Data Model	9
4. AI Assistant Module	10
4.1 NLP Chatbot Architecture	10
4.2 Knowledge Retrieval System	10
4.3 Multi-Turn Conversation	11
4.4 Smart Escalation	11
4.5 AIChatLog Data Model	12
5. AI-Powered Recruitment	12
5.1 Candidate Screening	13
5.2 Score-Based Ranking	13
5.3 Source Analytics	14
6. AI in Performance Management	14
6.1 Predictive Analytics	14

6.2 Skill Gap Identification	15
6.3 Recommendation Engine	15
7. AI Integration Architecture	16
7.1 API Routes	16
7.2 SDK Usage Patterns	16
8. AI Configuration Guide	17
8.1 Prompt Engineering Best Practices	17
8.2 Scoring Threshold Configuration	18
8.3 Model Selection Guide	18
9. AI Monitoring & Analytics	19
9.1 Usage Tracking	19
9.2 Cost Analysis	19
9.3 Performance Metrics	20
10. AI Ethics & Bias Mitigation	20
10.1 Ethical AI Principles	21
10.2 Bias Mitigation Strategies	21
10.3 AI Governance Framework	22

1. AI Platform Overview

The NEXUS HRMS platform integrates artificial intelligence capabilities throughout the employee lifecycle, powered by the z-ai-web-dev-sdk. AI features span recruitment screening, interview automation, conversational assistance, performance analytics, and predictive insights — transforming traditional HR processes into intelligent, data-driven workflows.

1.1 z-ai-web-dev-sdk Integration

The z-ai-web-dev-sdk serves as the foundational AI integration layer for the NEXUS platform. It provides standardized access to large language models for text generation, analysis, and conversational AI capabilities.

- **SDK Architecture:** The z-ai-web-dev-sdk is a JavaScript/TypeScript SDK that provides programmatic access to AI models through a unified API. It abstracts model selection, prompt engineering, and response parsing.
- **Integration Points:** The SDK is integrated into NEXUS through Next.js API routes, which serve as server-side AI orchestration points. Client-side components interact with AI features through these API endpoints.
- **Authentication:** The SDK authenticates using API keys stored as server-side environment variables, never exposed to the client.
- **Rate Management:** The SDK includes built-in rate limiting and retry logic to manage API call volumes and costs.
- **Error Handling:** Comprehensive error handling with fallback responses ensures graceful degradation when AI services are unavailable.
- **Response Streaming:** For chat interfaces, the SDK supports streaming responses to provide real-time feedback to users.

1.2 AI Capabilities Across the Platform

Module	AI Capability	SDK Integration	Business Value
AI Interview	Question generation, response evaluation, scoring	/api/ai-interview	Consistent, unbiased candidate assessment
AI Assistant	NLP chatbot, knowledge retrieval, escalation	/api/ai-chat	24/7 employee self-service, reduced HR workload
Recruitment	Resume screening, score-based ranking	/api/ai-interview	Faster hiring, better candidate matching

Performance	Predictive analytics, skill gap identification	/api/ai-chat	Proactive talent management
Helpdesk	Auto-categorization, smart routing	/api/ai-chat	Faster resolution, improved satisfaction
Training	Content recommendations, assessment generation	/api/ai-chat	Personalized learning paths

2. AI Admin Console

The AI Admin Console provides Super Admins and Tenant Admins with centralized control over AI features, model configuration, prompt templates, monitoring, and cost tracking. It serves as the governance layer for all AI operations within the platform.

2.1 Model Configuration

The Model Configuration section allows administrators to select and configure AI models for different use cases:

- **Model Selection:** Choose from available AI models (e.g., GPT-4, GPT-3.5, Claude) for each feature module. Different models may be optimal for different tasks.
- **Temperature Setting:** Configure the creativity/precision balance (0.0 = deterministic, 1.0 = creative). Interview question generation uses moderate temperature; scoring uses low temperature.
- **Max Tokens:** Set maximum response length per request type to control output length and costs.
- **Top-P (Nucleus Sampling):** Configure the probability threshold for token selection, affecting response diversity.
- **Frequency Penalty:** Adjust repetition avoidance in generated content.
- **Presence Penalty:** Control topic diversity in generated responses.
- **Fallback Configuration:** Set up fallback model chains in case the primary model is unavailable or rate-limited.
- **A/B Testing:** Route a percentage of requests to different model configurations for performance comparison.

2.2 Prompt Templates

Prompt templates define the system prompts and instruction sets that guide AI behavior across different features:

Template	Purpose	Key Variables	Model Type
Interview Questions	Generate role-specific interview questions	jobTitle, skills, difficulty, category	High-quality text generation
Response Evaluation	Score candidate interview responses	question, answer, rubric, maxScore	Analytical/scoring
Feedback Generation	Create detailed interview feedback	scores, responses, jobRequirements	Structured output
Bias Detection	Flag potential scoring bias	allScores, demographics, criteria	Classification
Chat Welcome	Initialize AI Assistant conversation	userName, role, department	Conversational
Knowledge Retrieval	Formulate HR policy responses	query, policyContext, employeeLevel	RAG-based
Screening Summary	Summarize candidate screening results	resumeData, jobRequirements, scores	Summarization
Performance Insights	Generate performance trend analysis	reviewData, goals, history	Analytical

2.3 Monitoring Dashboard

The AI Monitoring Dashboard provides real-time visibility into AI system performance and health:

- **Request Volume:** Track the number of AI API calls per hour/day/week, segmented by module and model.
- **Response Time:** Monitor average and percentile response times for AI operations, with alerting on degradation.
- **Success Rate:** Track the percentage of successful AI responses versus errors or timeouts.
- **Token Usage:** Monitor input and output token consumption per module for cost management.
- **Quality Metrics:** Track user satisfaction ratings for AI-generated content and suggestions.
- **Error Analysis:** Categorize and track error types (rate limit, timeout, content filter, invalid response) for troubleshooting.
- **Usage Patterns:** Visualize usage patterns across time of day, day of week, and module to optimize resource allocation.

- **Health Status:** Real-time health indicators for each AI service endpoint with automatic alerting on outages.

2.4 Cost Tracking & Budget Management

The cost tracking module provides comprehensive financial oversight of AI operations:

Feature	Description	Configuration
Per-Request Cost	Calculate cost for each AI API call based on token usage and model pricing	Model pricing table maintained by admin
Budget Alerts	Trigger notifications when spending approaches or exceeds defined thresholds	Monthly budget per module and overall
Cost Allocation	Attribute AI costs to specific modules, tenants, and user groups	Tenant-level and module-level tracking
Usage Forecasting	Project future AI costs based on usage trends	Historical data analysis with trend projection
Cost Optimization	Identify opportunities to reduce costs (model downgrades, caching)	Automatic suggestions based on usage patterns
Invoice Generation	Generate AI cost reports for internal chargeback	Monthly/weekly automated reports

3. AI Interview Module

The AI Interview Module is one of the most sophisticated AI features in NEXUS HRMS, providing automated interview question generation, response evaluation, scoring, feedback generation, and bias detection. It leverages the `z-ai-web-dev-sdk` to deliver consistent, fair, and comprehensive candidate assessments.

3.1 Question Generation

The AI interview question generation system creates role-specific, competency-based questions tailored to the job requirements:

- **Input Parameters:** Job title, required skills list, experience level, interview category (technical, behavioral, situational), difficulty level, and number of questions.
- **Question Types:** The system generates multiple question types including open-ended, scenario-based, problem-solving, behavioral (STAR method), and

technical knowledge questions.

- **Competency Mapping:** Questions are mapped to specific competencies (communication, leadership, technical skills, problem-solving) for structured evaluation.
- **Difficulty Calibration:** Question difficulty is calibrated using the configured difficulty parameter and the job's seniority level to ensure appropriate challenge.
- **Diversity Assurance:** The generation algorithm ensures question diversity across competency areas, avoiding repetition and ensuring comprehensive assessment coverage.
- **Custom Question Integration:** Recruiters can add custom questions to the AI-generated set, creating a hybrid interview that combines AI consistency with human expertise.
- **Question Validation:** Generated questions undergo automated validation for clarity, relevance, and bias before being presented to candidates.
- **Question Bank:** Successfully validated questions are stored in a question bank for reuse and continuous improvement of the generation model.

3.2 Automated Assessment

The automated assessment system evaluates candidate responses against predefined rubrics and competency frameworks:

- **Response Analysis:** AI analyzes the semantic content, structure, and completeness of candidate responses to interview questions.
- **Rubric-Based Scoring:** Each question has an associated scoring rubric that defines criteria for different score levels. The AI maps response content to rubric criteria.
- **Competency Assessment:** Responses are evaluated against the specific competency being measured (e.g., leadership, problem-solving, communication).
- **Contextual Understanding:** The AI considers the job context and requirements when evaluating response relevance and quality.
- **Multi-Factor Evaluation:** Scoring considers multiple factors including accuracy, depth, relevance, structure, and communication clarity.
- **Comparative Analysis:** When multiple candidates interview for the same role, the system provides comparative assessment across candidates.
- **Confidence Scoring:** Each automated assessment includes a confidence score indicating the AI's certainty in the evaluation, flagging low-confidence assessments for human review.

3.3 Scoring Algorithms

The scoring system uses a multi-dimensional algorithm that produces the aiScore field in the Interview model:

Scoring Dimension	Weight	Description	Scale
Technical Accuracy	30%	Correctness of technical knowledge demonstrated	0-10
Problem Solving	25%	Approach to problem-solving and analytical thinking	0-10
Communication	20%	Clarity, structure, and effectiveness of response	0-10
Relevance	15%	Direct relevance of response to the question asked	0-10
Depth	10%	Depth of understanding beyond surface-level answers	0-10

The composite aiScore is calculated as a weighted average of all dimensions, producing a score on a 0-10 scale. Thresholds for pass/fail are configurable per role through the AI Admin Console.

3.4 Feedback Generation

The feedback generation system creates comprehensive, actionable feedback stored in the aiFeedback field of the Interview model:

- **Strengths Summary:** Identifies and highlights the candidate's strongest competency areas with specific examples from their responses.
- **Improvement Areas:** Identifies areas where the candidate fell short, with constructive suggestions for improvement.
- **Competency Breakdown:** Provides detailed scores for each competency assessed, enabling targeted development planning.
- **Response-Level Feedback:** Offers specific feedback on individual responses, referencing what was said and how it could be improved.
- **Comparative Context:** Where applicable, provides context on how the candidate's performance compares to typical scores for the role level.
- **Recommendation:** Generates a hiring recommendation (Strong Hire, Hire, Lean Hire, No Hire) based on the composite score and threshold configuration.

- **Interview Summary:** Produces a concise executive summary suitable for sharing with the hiring committee.

3.5 Bias Detection

The bias detection module identifies potential biases in AI scoring to ensure fair and equitable candidate assessment:

- **Score Distribution Analysis:** Monitors the distribution of AI scores across demographic groups to identify statistically significant disparities.
- **Question Bias Detection:** Analyzes whether specific questions produce systematically different scores for certain demographic groups.
- **Language Bias Checks:** Evaluates whether the AI scoring model exhibits bias toward certain communication styles or language patterns.
- **Flagging System:** When potential bias is detected, the system flags the interview for mandatory human review before any hiring decisions are made.
- **Bias Report Generation:** Generates periodic bias analysis reports for the AI Admin Console, highlighting trends and areas requiring model recalibration.
- **Mitigation Strategies:** The system suggests mitigation strategies when bias patterns are identified, such as adjusting prompt templates, adding debiasing instructions, or modifying scoring weights.

3.6 Interview Data Model

The Interview model in the NEXUS database stores both traditional and AI-generated assessment data:

Field	Type	Description
id	Int (Auto)	Unique interview identifier
candidateId	String	Reference to the candidate/employee
jobId	String	Reference to the job posting
interviewDate	DateTime	When the interview was conducted
interviewType	String	AI / Manual / Hybrid
questions	Json	Array of interview questions with metadata
responses	Json	Array of candidate responses
aiScore	Float (Optional)	Composite AI-generated score (0-10)
aiFeedback	String (Optional)	AI-generated feedback text
manualScore	Float (Optional)	Human evaluator score

manualFeedback	String (Optional)	Human evaluator feedback
biasFlags	Json (Optional)	Array of detected bias indicators
status	String	Scheduled / In Progress / Completed / Reviewed
tenantId	String	Tenant isolation foreign key

4. AI Assistant Module

The AI Assistant Module provides an intelligent chatbot interface for employee self-service and HR query resolution. Powered by the z-ai-web-dev-sdk, the assistant understands natural language queries, retrieves relevant HR knowledge, and performs actions on behalf of the employee.

4.1 NLP Chatbot Architecture

The AI Assistant uses a multi-stage NLP pipeline to understand and respond to employee queries:

- **Intent Classification:** The first stage identifies the user's intent (e.g., leave inquiry, policy question, helpdesk request, payroll question) using the AI model with a predefined intent taxonomy.
- **Entity Extraction:** Key entities are extracted from the query (e.g., dates, leave types, policy names, amounts) to provide context for the response.
- **Context Resolution:** The system resolves references using conversation context and the user's profile data (e.g., "my leave balance" resolves to the specific employee's leave records).
- **Knowledge Retrieval:** Based on the classified intent and extracted entities, the system retrieves relevant knowledge from the HR knowledge base and policy documents.
- **Response Generation:** The AI model generates a contextual, accurate response incorporating retrieved knowledge, user context, and conversational history.
- **Action Execution:** For actionable requests (e.g., "submit a leave request"), the assistant can execute API actions on behalf of the user with confirmation.

4.2 Knowledge Retrieval System

The knowledge retrieval system ensures the AI Assistant provides accurate, organization-specific responses:

- **Policy Knowledge Base:** Organizational policies (leave, attendance, travel, expenses) are indexed and made searchable by the AI system.
- **FAQ Repository:** Common questions and their approved answers are maintained in a structured FAQ database.
- **Process Documentation:** Step-by-step process documentation for common HR tasks is available for the AI to reference.
- **Benefit Information:** Employee benefit details, eligibility criteria, and claim procedures are indexed for retrieval.
- **Real-Time Data:** For queries about personal data (leave balance, payslip, attendance), the system queries live data with appropriate access controls.
- **Context-Aware Retrieval:** The retrieval system considers the user's role, department, and location to provide contextually relevant information.

4.3 Multi-Turn Conversation

The AI Assistant maintains conversation context across multiple messages, enabling natural multi-turn dialogues:

- **Session Management:** Each conversation session maintains a unique session ID with full message history for context continuity.
- **Context Window:** The system maintains a configurable context window of recent messages to inform the AI model about the ongoing conversation.
- **Reference Resolution:** Pronouns and references to previous messages are resolved using conversation context (e.g., "What about for next month?" after a leave balance query).
- **Topic Tracking:** The system tracks the current topic of conversation and can detect topic shifts, adapting responses accordingly.
- **Clarification Requests:** When user queries are ambiguous, the assistant asks clarifying questions rather than making assumptions.
- **Conversation Summarization:** For long conversations, the system summarizes previous context to maintain coherence without exceeding token limits.

4.4 Smart Escalation

When the AI Assistant encounters queries beyond its capability, it intelligently escalates to human agents:

- **Capability Threshold:** Queries requiring subjective judgment, exception handling, or sensitive decisions trigger automatic escalation.

- **Confidence-Based Escalation:** When the AI's response confidence falls below a configurable threshold, the conversation is escalated.
- **Emotion Detection:** The system detects frustrated or distressed language patterns and prioritizes escalation to human support.
- **Seamless Handoff:** Escalated conversations include full chat history and context, enabling the human agent to continue without requiring the user to repeat information.
- **Category-Based Routing:** Escalated queries are routed to the appropriate team (HR, IT, Finance, Helpdesk) based on the classified intent.
- **Priority Assignment:** Escalation priority is determined by the query category, user role, and detected urgency level.

4.5 AIChatLog Data Model

All AI Assistant interactions are recorded in the AIChatLog model for analytics and improvement:

Field	Type	Description
id	Int (Auto)	Unique chat log identifier
userId	String	User who initiated the conversation
sessionId	String	Conversation session identifier
message	String	User message content
response	String	AI assistant response
intent	String (Optional)	Classified intent category
entities	Json (Optional)	Extracted entities from message
confidence	Float (Optional)	AI response confidence score
escalated	Boolean	Whether the conversation was escalated
escalatedTo	String (Optional)	Agent/team the conversation was escalated to
tokensUsed	Int (Optional)	Total tokens consumed for this interaction
responseTime	Int (Optional)	AI response time in milliseconds
tenantId	String	Tenant isolation foreign key
timestamp	DateTime	When the interaction occurred

5. AI-Powered Recruitment

AI integration in the recruitment module enhances every stage of the hiring pipeline from candidate screening to final selection. The AI capabilities reduce time-to-hire, improve candidate quality, and minimize unconscious bias in the selection process.

5.1 Candidate Screening

AI-powered candidate screening analyzes resumes and applications against job requirements to produce a relevance score:

- **Resume Parsing:** AI extracts structured data from resumes including skills, experience, education, certifications, and projects.
- **Skills Matching:** The system matches extracted skills against job requirements, accounting for skill synonyms and related competencies.
- **Experience Assessment:** AI evaluates the depth and relevance of work experience in relation to the target role.
- **Education Verification:** Educational qualifications are matched against job requirements with consideration for equivalencies.
- **Cultural Fit Indicators:** The system analyzes resume language and project descriptions for alignment with organizational culture markers.
- **Gap Identification:** AI identifies gaps between the candidate profile and job requirements, highlighting areas for interview focus.
- **Scoring Output:** Each candidate receives a composite screening score (0-100) that ranks them against other applicants for the same position.

5.2 Score-Based Ranking

The score-based ranking system provides an objective, data-driven method for prioritizing candidates:

Score Range	Category	Recommended Action	Typical Percentage
90-100	Exceptional Match	Fast-track interview scheduling	5-10%
75-89	Strong Match	Priority interview scheduling	15-25%
60-74	Good Match	Standard interview queue	25-35%
45-59	Partial Match	Review for alternative roles or defer	20-30%
0-44	Weak Match	Politely decline or archive	15-25%

5.3 Source Analytics

AI-powered source analytics track the effectiveness of different recruitment channels:

- **Channel Performance:** Track conversion rates from each sourcing channel (job boards, referrals, social media, career site) through the entire hiring funnel.
- **Quality by Source:** Analyze the average candidate quality (as measured by AI screening scores) from each source to optimize recruitment spend.
- **Cost-Per-Hire Analysis:** Calculate the cost-per-hire for each sourcing channel, enabling data-driven budget allocation.
- **Time-to-Hire by Source:** Track how quickly candidates from different sources move through the hiring pipeline.
- **Retention Correlation:** Analyze the correlation between sourcing channel and employee retention to identify channels that produce longer-tenured employees.
- **Predictive Sourcing:** AI recommends optimal sourcing channel mix based on historical performance data for similar roles.

6. AI in Performance Management

AI integration in the Performance Management module provides predictive analytics, skill gap identification, and intelligent recommendations that transform performance management from a reactive, periodic process to a proactive, continuous improvement system.

6.1 Predictive Analytics

The predictive analytics engine uses historical performance data to forecast future outcomes and identify at-risk employees:

- **Flight Risk Prediction:** AI analyzes patterns in attendance, engagement, performance trends, and compensation data to identify employees at risk of leaving the organization.
- **Performance Trajectory:** Models predict future performance ratings based on historical trends, enabling proactive intervention for declining performers.
- **Promotion Readiness:** AI assesses an employee's readiness for promotion based on performance history, skill development, and role requirements alignment.

- **Attrition Impact Analysis:** Predicts the organizational impact of potential employee departures, enabling succession planning.
- **Team Performance Forecasting:** Projects team-level performance based on composition, workload, and historical patterns.
- **Calibration Assistance:** AI identifies potential rating inconsistencies across managers, suggesting calibration adjustments for fairness.

6.2 Skill Gap Identification

The AI skill gap analysis identifies discrepancies between current employee capabilities and organizational needs:

- **Current Skill Mapping:** AI analyzes employee performance data, training completions, and project histories to build a comprehensive skill profile for each employee.
- **Required Skill Mapping:** Job role requirements and organizational competency frameworks define the target skill set for each position.
- **Gap Calculation:** The system calculates the delta between current and required skills, quantifying gaps by severity and business impact.
- **Team-Level Aggregation:** Skill gaps are aggregated at the team and department levels to identify systemic capability deficiencies.
- **Priority Ranking:** Gaps are ranked by business impact (critical skills missing for key roles) and feasibility (training time and cost to close the gap).
- **Development Pathways:** AI recommends specific training programs, mentorship assignments, and project experiences to close identified skill gaps.

6.3 Recommendation Engine

The recommendation engine generates personalized suggestions for employee development and organizational improvement:

Recommendation Type	Inputs	Output	Frequency
Training Recommendations	Skill gaps, career goals, available courses	Personalized learning path	Monthly / On-demand
Career Path Suggestions	Skills, aspirations, organizational needs	Suggested career trajectory	Quarterly
Mentorship Matching	Skill gaps, available mentors, compatibility	Mentor-mentee pairing	As needed
Project Assignment	Skills, development goals, project needs	Optimal project assignment	Per project

Performance Intervention	Declining trends, engagement data	Intervention recommendation	Real-time alerts
Succession Planning	Role criticality, employee readiness	Succession candidates ranked	Semi-annually

7. AI Integration Architecture

This section details the technical architecture of AI integration in the NEXUS platform, covering API routes, SDK usage patterns, and data flow between components.

7.1 API Routes

AI functionality is exposed through dedicated Next.js API routes that serve as server-side AI orchestration endpoints:

Route	Method	Purpose	Authentication
/api/ai-interview	POST	Generate questions, evaluate responses, produce scores	JWT Bearer Token
/api/ai-interview	GET	Retrieve interview results and AI feedback	JWT Bearer Token
/api/ai-chat	POST	Send message to AI Assistant, receive response	JWT Bearer Token
/api/ai-chat	GET	Retrieve chat history for a session	JWT Bearer Token
/api/ai-admin/config	GET/PUT	Manage AI model configurations	JWT Bearer (Admin only)
/api/ai-admin/templates	GET/POST/PUT	Manage prompt templates	JWT Bearer (Admin only)
/api/ai-admin/analytics	GET	Retrieve AI usage and performance metrics	JWT Bearer (Admin only)

7.2 SDK Usage Patterns

The `z-ai-web-dev-sdk` is used in NEXUS following consistent patterns for reliability and maintainability:

- **Server-Side Only:** All SDK calls are made from server-side API routes, never from client components. This protects API keys and allows server-side rate management.

- **Request Validation:** All AI API routes validate the request payload, user authentication, and authorization before making SDK calls.
- **Error Handling:** SDK calls are wrapped in try-catch blocks with graceful fallback responses when AI services are unavailable.
- **Token Tracking:** Each SDK call tracks input and output token counts for cost management and usage analytics.
- **Response Validation:** AI responses are validated for format and content before being returned to the client or stored in the database.
- **Caching Strategy:** Frequently requested AI-generated content (e.g., common FAQ responses) is cached to reduce SDK calls and costs.
- **Retry Logic:** The SDK implements exponential backoff retry logic for transient failures, with a maximum retry count to prevent indefinite retries.
- **Timeout Configuration:** Each SDK call has a configured timeout to prevent long-running requests from blocking the application.

8. AI Configuration Guide

This section provides practical guidance for configuring AI features in NEXUS HRMS, including prompt engineering best practices, scoring threshold configuration, and model selection strategies.

8.1 Prompt Engineering Best Practices

- **Be Specific and Structured:** Provide clear, specific instructions in system prompts. Use structured output formats (JSON) for AI responses that need to be parsed programmatically.
- **Include Context:** Always include relevant context in prompts — job details for interview questions, policy text for assistant responses, scoring rubrics for evaluations.
- **Define Output Format:** Explicitly specify the expected output format, structure, and length. Example: "Generate exactly 5 questions in JSON format with fields: question, category, difficulty, competency."
- **Set Role and Persona:** Define a clear role for the AI in the system prompt. Example: "You are an experienced HR interviewer evaluating candidates for a senior software engineer position."
- **Use Examples (Few-Shot):** Include 1-2 examples of desired input-output pairs in the prompt to guide the AI toward the expected response style and quality.

- **Add Constraints:** Specify what the AI should NOT do. Example: "Do not generate questions about salary expectations. Do not include any biased language regarding age, gender, or ethnicity."
- **Version Your Prompts:** Maintain versioned prompt templates in the AI Admin Console with change logs. This enables A/B testing and rollback if a new prompt version underperforms.
- **Test Iteratively:** Test prompts with diverse inputs before deployment. Evaluate output quality, consistency, and bias across multiple runs with different inputs.

8.2 Scoring Threshold Configuration

Scoring thresholds determine how AI-generated scores are interpreted and what actions are triggered:

Configuration	Default	Range	Description
Interview Pass Threshold	6.0	0.0 - 10.0	Minimum aiScore for candidate to pass AI interview
Auto-Shortlist Threshold	8.0	0.0 - 10.0	Score above which candidate is auto-shortlisted
Auto-Reject Threshold	3.0	0.0 - 10.0	Score below which candidate is auto-rejected (with review)
Bias Alert Threshold	1.5 std dev	0.5 - 3.0	Standard deviation threshold for bias flagging
Confidence Threshold	0.7	0.0 - 1.0	Minimum AI confidence for auto-processing; below triggers human review
Escalation Confidence	0.4	0.0 - 1.0	Confidence below which AI Assistant escalates to human
Screening Score Weight	0.4	0.0 - 1.0	Weight of AI screening score in composite candidate ranking

8.3 Model Selection Guide

Selecting the appropriate AI model for each use case balances quality, cost, and latency:

Use Case	Recommended Model	Rationale	Cost Level
Interview Question Generation	GPT-4 / Claude Opus	High quality, nuanced questions	High

Response Evaluation & Scoring	GPT-4 / Claude Sonnet	Accurate, consistent scoring	Medium-High
Feedback Generation	GPT-4 / Claude Sonnet	Detailed, constructive feedback	Medium-High
Chat Assistant (General)	GPT-3.5 / Claude Haiku	Fast, cost-effective for FAQ	Low
Chat Assistant (Complex)	GPT-4 / Claude Sonnet	Complex queries need better reasoning	Medium-High
Resume Screening	GPT-3.5 / Claude Haiku	Structured extraction, bulk processing	Low-Medium
Performance Analytics	GPT-4 / Claude Sonnet	Nuanced analysis and recommendations	Medium-High
Bias Detection	GPT-4 / Claude Opus	Requires strong reasoning for fairness	High

9. AI Monitoring & Analytics

Comprehensive monitoring and analytics are essential for maintaining AI system quality, managing costs, and ensuring ethical AI deployment. This section covers the monitoring framework, key metrics, and analytics capabilities.

9.1 Usage Tracking

Metric	Granularity	Collection Method	Retention
API Call Count	Per route, per user, per tenant	Request logging middleware	12 months
Token Consumption	Per call, per module	SDK response metadata	12 months
Model Distribution	Per module, per model type	Configuration + request logging	12 months
User Adoption	Per module, per role	Feature usage tracking	Ongoing
Session Metrics	Per chat session, per interview	Session start/end logging	6 months
Escalation Rate	Per module, per reason	Escalation event logging	12 months

9.2 Cost Analysis

Cost analysis provides financial visibility into AI operations:

- **Per-Request Cost Calculation:** Each AI API call cost = (input tokens × input price) + (output tokens × output price) based on the model's published pricing.
- **Monthly Cost Trends:** Track month-over-month AI spending by module, model, and tenant to identify cost growth patterns.
- **Cost Per Outcome:** Calculate the cost of AI operations per business outcome (e.g., cost per interview conducted, cost per chat session resolved).
- **Budget Utilization:** Monitor budget consumption against allocated limits with alerting at 50%, 75%, 90%, and 100% thresholds.
- **Optimization Opportunities:** AI-driven identification of cost-saving opportunities such as model downgrades for simple tasks, caching for repeated queries, and batch processing for bulk operations.
- **ROI Estimation:** Estimate the return on AI investment by comparing AI operational costs against time saved, quality improvements, and error reduction metrics.

9.3 Performance Metrics

Metric	Target	Alert Threshold	Measurement
Average Response Time	< 3 seconds	> 5 seconds	P50 latency per API route
P95 Response Time	< 8 seconds	> 12 seconds	P95 latency per API route
Success Rate	> 99%	< 97%	Successful responses / total requests
Timeout Rate	< 0.1%	> 0.5%	Timeout errors / total requests
Content Filter Rate	< 1%	> 3%	Filtered responses / total requests
User Satisfaction	> 4.0/5	< 3.5/5	User rating of AI interactions
First-Contact Resolution	> 70%	< 50%	AI-resolved / total chat sessions
Escalation Rate	< 20%	> 35%	Escalated sessions / total sessions

10. AI Ethics & Bias Mitigation

Ethical AI deployment is paramount in an HRMS platform where AI decisions directly impact people's careers and livelihoods. NEXUS implements a comprehensive ethics framework and bias mitigation strategy to ensure fair, transparent, and accountable AI operations.

10.1 Ethical AI Principles

- **Fairness:** AI systems must treat all individuals equitably regardless of protected characteristics (age, gender, race, religion, disability, etc.). All AI models undergo bias testing before deployment.
- **Transparency:** AI decision-making processes must be explainable. Users affected by AI decisions have the right to understand how the decision was made and what factors were considered.
- **Accountability:** Clear accountability structures must exist for AI decisions. No fully automated decision should be final without human review for high-impact outcomes (hiring, termination, compensation).
- **Privacy:** AI systems must respect data privacy principles, processing only necessary data and maintaining confidentiality. Personal data used for AI training must be anonymized.
- **Safety:** AI systems must be designed to prevent harm. Fail-safe mechanisms ensure that AI errors do not result in adverse outcomes for employees or candidates.
- **Human Oversight:** Critical AI decisions require human oversight. The system is designed as human-in-the-loop, not fully autonomous, for decisions affecting employment outcomes.

10.2 Bias Mitigation Strategies

- **Debiased Training Data:** Ensure training data represents diverse populations. Audit training datasets for representation gaps and augment underrepresented groups.
- **Blind Processing:** Remove personally identifiable information (names, photos, demographics) from AI input where possible to prevent demographic-based scoring variations.
- **Regular Bias Audits:** Conduct quarterly bias audits of all AI models, analyzing score distributions across demographic groups for statistically significant disparities.
- **Diverse Testing:** Test AI systems with diverse input data representing different demographics, communication styles, and backgrounds before deployment.
- **Adversarial Testing:** Intentionally test AI systems with inputs designed to trigger biased outputs, and use findings to improve model behavior.
- **Human Review Requirement:** For interview scoring and candidate ranking, AI results are advisory — final decisions must involve human review. The

system enforces this through workflow configuration.

- **Feedback Loop:** Implement a feedback mechanism where users can report perceived bias. All reports are investigated and findings inform model improvements.
- **Transparency Reports:** Publish periodic AI fairness reports showing score distributions, bias detection results, and mitigation actions taken.

10.3 AI Governance Framework

The AI Governance Framework establishes organizational structures and processes for responsible AI deployment:

Component	Description	Responsible Party
AI Ethics Committee	Reviews and approves AI feature deployments, audits bias reports	Leadership + HR + Legal
Model Approval Process	New models and prompt templates require review before production use	AI Admin + Ethics Committee
Incident Response	Process for handling AI-related incidents (bias reports, errors)	HR + AI Admin
Regular Audits	Quarterly bias and performance audits of all AI features	AI Ethics Committee
Transparency Obligations	Requirements for disclosing AI usage to candidates and employees	Legal + HR
Continuous Monitoring	Ongoing monitoring of AI fairness metrics and performance	AI Admin + Engineering